

[Products](#) ▾[Cloud](#) ▾[Solutions](#) ▾[Developer?](#) ▾[Customers](#)[Pricing](#)[Blog](#)[Contact](#)[Login](#) ▾[Create Account](#)

# Postgres Internals Hiding in Plain Sight

Postgres has an awesome amount of data collected in its own internal tables. Postgres hackers know all about this - but software developers and folks working with day to day Postgres tasks often miss out the good stuff.

The Postgres catalog is how Postgres keeps track of itself. Of course, Postgres would do this in a relational database with its own schema. Throughout the years several nice features have been added to the internal tables like psql tools and views that make navigating Postgres' internal tables even easier.

Today I want to walk through some of the most important Postgres internal data catalog details. What they are, what is in them, and how they might help you understand more about what is happening inside your database.

## psql's catalog information

The easiest way to get at some of Postgres' internal catalogs is to use the built-in [psql commands](#) that begin \d generally. Here's some common Postgres ones users should be comfortable using:

**\d {tablename}** : describes a specific table. \d will do a lot of things if you qualify \d with a table or view name.

**\di** : list all your indexes

**\dx** : list installed extensions

**\dp** : to show access privileges

**\dp+** : tables and views with the roles and access details

**\dconfig** : your current configuration settings

**\dt {tablename}** : describe a table

**\dti+** : tables and indexes with sizes

**\dg+** : show role names

**\df** : show your functions

**\dv {view name}** : describe a view

**\l** : lists all your databases

## Important Postgres catalog views

Postgres exposes many of the complex internals of the database system in easy-to-query views. These host a wealth of information about what is going on inside your database and direct SQL access to answer in the

moment emergency questions like “what is taking up all my CPU” and more long term questions like “what are my 10 slowest queries”.

## pg\_stat\_activity

Shows current database activity, including running queries, state, and client information. Essential for troubleshooting and getting process ids (pid) for bad actors.

```
SELECT pid, username, datname, client_addr, application_name, state, query
FROM pg_stat_activity
WHERE state != 'idle'
ORDER BY state, query_start DESC;
```

## pg\_stat\_statements

This requires the pg\_stat\_statements extension - but it is part of the contrib library and ships with Postgres, so doesn't require separate installation.

This view tracks execution statistics for all queries executed by all databases. It's incredibly powerful for identifying slow or frequently executed queries.

```
-- pg_stat_statements 10 longest running queries
SELECT query, calls, total_exec_time, mean_exec_time, rows
FROM pg_stat_statements
ORDER BY total_exec_time DESC
LIMIT 10;
```

## pg\_stat\_database

This view provides database-wide statistics, such as the number of connections, transactions, and I/O. It's useful for a high-level overview of database activity and health.

```
-- high level db stats for the postgres db
SELECT datname, numbackends, xact_commit, xact_rollback, blks_read, blks_hit
FROM pg_stat_database
WHERE datname = 'postgres';
```

## pg\_locks

This view displays information about locks held by active processes. This is the go to place for troubleshooting locking issues, deadlocks, and contention within the database. We have a great blog on locking and [how to find the source of the lock in Postgres](#).

```
-- Locks joined with the activity table. Shows not granted Locks, typically those that could not be granted because they
SELECT a.datname, l.pid, l.locktype, l.relation::regclass, l.mode, l.granted
FROM pg_locks l
JOIN pg_stat_activity a ON l.pid = a.pid
WHERE NOT l.granted;
```

## pg\_stat\_user\_tables

This view offers statistics on tables, including sequential scans, index scans, and row-level operations (inserts, updates, deletes). It's great for identifying tables with heavy activity or those that need vacuuming.

```
-- see sequence scans and index scans by table
SELECT relname AS table_name, seq_scan, idx_scan
FROM pg_stat_user_tables
WHERE seq_scan > 0 OR idx_scan > 0 ORDER BY seq_scan DESC;
```

## pg\_stat\_user\_indexes

This view provides statistics on user indexes, such as how often they're used and how many tuples are read. This is particularly helpful for finding unused or underutilized indexes.

```
-- Never used indexes in Postgres sorted by size
SELECT s.schemaname, s.relname AS table_name, s.indexrelname AS index_name, pg_size.pretty(pg_relation_size(s.indexrel
FROM pg_stat_user_indexes AS s
JOIN pg_index AS i ON s.indexreloid = i.indexreloid
WHERE s.idx_scan = 0 AND i.indisunique IS FALSE
ORDER BY pg_relation_size(s.indexreloid) DESC;
```

## pg\_settings

This is a prebuilt view that is super useful for viewing configuration parameters, their current values, and their descriptions. Qualify with **ILIKE** to see exact parameters you're looking for.

```
-- find shared_buffer or work_mem settings
SELECT name, setting, unit, short_desc
FROM pg_settings
WHERE name LIKE '%shared_buffers%' OR name LIKE '%work_mem%';
```

## pg\_roles

This view describes all system roles, which include users and groups. It's useful for checking permissions, login capabilities, and role memberships.

```
-- This query lists all roles, showing their names, whether they can log in, and their password expiration date.  
SELECT rolname, rolcanlogin, rolvaliduntil  
FROM pg_roles  
ORDER BY rolname;
```

## pg\_database

This view contains all databases in the cluster. It provides key metadata for each database, including its owner, character encoding, and access privileges. We have a lot of folks now that create dozens and sometimes hundreds of databases for development, so this is a good high level view.

```
-- This query lists all Postgres databases, their sizes, and owners.  
SELECT d.datname AS database_name, pg_size_pretty(pg_database_size(d.datname)) AS database_size, pg_get_userbyid(d.dat  
FROM pg_database AS d  
WHERE d.datistemplate = false;
```

## Postgres catalog tables

Behind the Postgres metacommands and views - there are several core catalog tables. Many of the psql commands match up with the catalog tables. Something roughly like this:

psql command	what data	catalog tables
\d	tables and table objects	pg_class
\di	indexes	pg_class, pg_index
\dx	installed extensions	pg_extension
\dp	tables and privileges	pg_class, pg_roles, pg_attribute
\l	databases	pg_database
\df	available functions	pg_proc

Let's look at these and how you might want to use them.

## pg\_stats

The pg\_stats table collects all the details about your columns - things like cardinality - are there many items in this column or a few? Postgres uses a lot of the details in pg\_stats to make decisions for the query planner and efficiently. In some cases, giving [pg\\_stats more information can make your queries faster.](#)

```
-- table column data Like cardinality
SELECT * FROM pg_stats
WHERE tablename = 'table_name'
AND attname = 'column_name';
```

## pg\_class

`pg_class` contains a row for every table, index, sequence, view, materialized view, and other "relation-like" objects in the database. Sometimes this is a nice high level view of an entire table's accoutrements.

```
SELECT c.relname, pg_get_userbyid(c.relpowner) AS owner
FROM pg_class c
JOIN pg_namespace n ON n.oid = c.relnamespace
WHERE n.nspname = 'public' AND c.relkind = 'r'
ORDER BY c.relname;
```

## pg\_type

This table stores all data types that exist. It's confusing though - in Postgres, every table has an associated composite type that defines the structure of its rows. So if you do a `select *` you'll see all the table names here and all the data types. If you filter a bit, you can see all your custom data types, domains, and enums.

```
-- see your custom data types in Postgres
SELECT
    t.typname AS type_name,
    n.nspname AS schema_name,
    t.typtype AS type_class
FROM
    pg_type AS t
JOIN
    pg_namespace AS n ON t.typnamespace = n.oid
LEFT JOIN pg_class c ON typrelid = c.oid
WHERE
    t.typtype IN ('e', 'd', 'c') -- 'e' for enum, 'd' for domain, 'c' for composite types.
    AND n.nspname NOT IN ('pg_catalog', 'information_schema', 'pg_toast')
    AND (t.typtype <> 'c' OR c.relkind = 'c')
```

```
ORDER BY  
schema_name, type_name;
```

## pg\_proc

This is the catalog of all functions and stored procedures that Postgres can use. It contains metadata about each routine. Made a function last week but can't find it now? Just scan through all of them.

```
-- This query finds all functions, triggers, and stored procedures.  
SELECT proname AS function_name, proargnames AS argument_names, pg_catalog.format_type(prorettype, NULL) AS return_type  
FROM pg_proc  
ORDER BY proname;
```

## pg\_attribute

This table stores information about table columns and there is one row in **pg\_attribute** for every column in every table. While indexes and other objects that have an entry in **pg\_class**.

Query columns and data types for any table with a query like this:

```
SELECT  
a.attname AS column_name,  
pg_catalog.format_type(a.atttypid, a.atttypmod) AS data_type  
FROM  
pg_catalog.pg_attribute a  
WHERE  
a.attrelid = 'orders'::regclass  
AND a.attnum > 0
```

```
AND NOT a.attisdropped  
ORDER BY  
    a.attnum;
```



## pg\_catalog schema

The pg\_catalog is the schema holding the system tables, so you will either need to include `pg_catalog` in your `search_path` (the default), or any query you issue will need to be qualified with `pg_catalog`.

Here's a summary of the internal catalog tables:

pg_catalog	schema holding all the catalog tables
pg_stats	table and column statistics, like cardinality
pg_attribute	row for every table column
pg_class	every table, index, view, materialized view, foreign table
pg_type	data types, built in and custom

Elizabeth Christensen

Nov 7, 2025 • 9 min read

More by this author

### LATEST ARTICLES

[Postgres Internals Hiding in Plain Sight](#)

[PostGIS Performance: Improve Bounding Boxes with Decompose and Subdivide](#)

[Temporal Joins](#)

[PostGIS Performance: pg\\_stat\\_statements and Postgres tuning](#)

[Is Postgres Read Heavy or Write Heavy? \(And Why You Should Care\)](#)

## Exploring system tables with ECHO\_HIDDEN or -E

Sometimes navigating these tables and views can be confusing and require browsing through a mix of docs and source code. If you want to have some fun exploring how the catalog is connected, you can connect to your database with `-E` argument to psql (or do `\set ECHO_HIDDEN` on if you're already connected). Postgres will echo each psql the command that's run with SQL so you can grab the underlying SQL and edit from there.

For example, echoing `\dt+` will show me a query and the results.

```
SELECT n.nspname as "Schema",
       c.relname as "Name",
      CASE c.relkind WHEN 'r' THEN 'table' WHEN 'v' THEN 'view' WHEN 'm' THEN 'materialized view' WHEN 'i' THEN 'index' WHEN 'f' THEN 'function' WHEN 'e' THEN 'event trigger' ELSE 'other' END as "Type",
      pg_catalog.pg_get_userbyid(c.relpersistence) as "Owner",
      CASE c.relpersistence WHEN 'p' THEN 'permanent' WHEN 't' THEN 'temporary' WHEN 'u' THEN 'unlogged' END as "Persistence",
      am.amname as "Access method",
      pg_catalog.pg_size.pretty(pg_catalog.pg_table_size(c.oid)) as "Size",
      pg_catalog.obj_description(c.oid, 'pg_class') as "Description"
  FROM pg_catalog.pg_class c
  LEFT JOIN pg_catalog.pg_namespace n ON n.oid = c.relnamespace
  LEFT JOIN pg_catalog.pg_am am ON am.oid = c.relam
 WHERE c.relkind IN ('r','p','')
    AND n.nspname <> 'pg_catalog'
    AND n.nspname !~ '^pg_toast'
    AND n.nspname <> 'information_schema'
    AND pg_catalog.pg_table_is_visible(c.oid)
 ORDER BY 1,2;
```

List of tables

-[ RECORD 1 ]-----

Schema	public
Name	articles
Type	table
Owner	dba
Persistence	permanent
Access method	heap
Size	16 kB
Description	

## Getting to Postgres internals

1. The easiest way to see internals is to start with the psql `\d` commands
2. The prebuilt views like `pg_stat_activity`, `pg_stat_statements`, `pg_locks`, and `pg_stat_user_indexes` are ready to go for easy querying and searching.
3. Going a step deeper, you can access the underlying internal Postgres tables, housed in the `pg_catalog` schema. `-E echo_hidden` can help you see the tables involved if you echo psql commands.

Enjoy this article?

You will love our newsletter!

Enter your email

Join The List

[Crunchy Postgres](#)  
[Crunchy Postgres for Kubernetes](#)  
[Crunchy Bridge](#)  
[Crunchy Certified PostgreSQL](#)  
[Crunchy PostgreSQL for Cloud Fou...](#)  
[Crunchy MLS PostgreSQL](#)  
[Crunchy Spatial](#)

[Enterprise PostgreSQL Sup...](#)  
[Ansible](#)  
[Red Hat Partner](#)  
[Trusted PostgreSQL](#)  
[Crunchy Data Subscription](#)  
  
**Migrate to Crunchy Data**  
[Migrate from RDS](#)  
[Migrate from Heroku](#)

[Customer Portal](#)  
[Documentation](#)  
[Postgres Tutorials](#)  
[Crunchy Bridge Walkthrough](#)  
[Postgres Operator Walkthro...](#)  
[Blog](#)  
[Events](#)

[About](#)  
[Team](#)  
[News](#)  
[Contact Us](#)  
[Newsletter](#)  
[Branding](#)  
[Privacy Notice – Recently Upda...](#)  
[Security](#)  
[Cookies Settings](#)

**Subscribe to the Crunchy Data Newsletter**  
and receive Postgres content every month.

Enter your email

Subscribe

© 2018-2025 Crunchy Data Solutions, Inc.

