17
Where To Go For More Learning

At the beginning of this book, we mentioned how PostgreSQL is a mature project with decades of development under its belt, and as would be expected, the software, community, and ecosystem around PostgreSQL are quite large; far too large for any one resource to cover in its entirety. The purpose of this book is to act as a starting point for learning PostgreSQL. To master PostgreSQL, more information will be needed. This chapter guides you on some of the best resources out there for learning PostgreSQL, and, how best to take advantage of those resources. In this chapter we'll cover:

- PostgreSQL documentation
- Books that help you learn PostgreSQL
- Events that include PostgreSQL
- Online PostgreSQL resources

As with the rest of the book, the information in this chapter is an introduction to what's out there, not a complete listing of all available resources. The most important aspect of using the resources listed here is to ensure that they are either the up-to-date version or the correct version for your PostgreSQL servers.

PostgreSQL Documentation

We reference the PostgreSQL documentation multiple times throughout the book. The reason we do is because that is your primary resource for official information about PostgreSQL. We always start most of our learning with the PostgreSQL documentation. Not only is it a useful resource, but, like PostgreSQL, the documentation is open source, so you can even contribute to making it better.

While the PostgreSQL documentation is your starting point, there are a few things that you have to keep an eye on. If we start with the root location for the documentation (no tinyurl for this one example), https://www.postgresql.org/docs, you'll see the following (as of this writing):

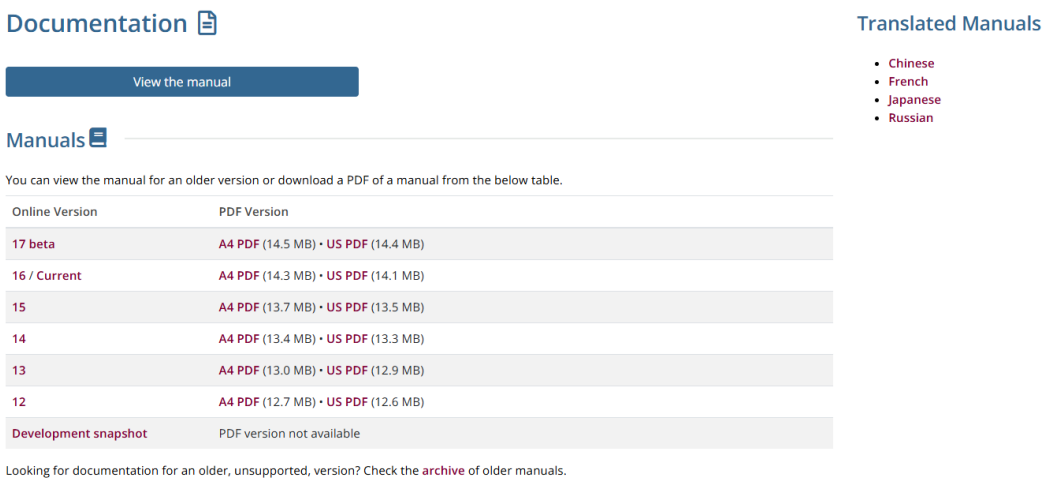

Figure 17-1. *PostgreSQL documentation landing page*

Two important points to note. First, you have various versions of PostgreSQL mentioned. The most important one, as of this writing, is the one for "16 / Current". The PostgreSQL documentation is designed so that "Current" always refers to the latest stable major version of Postgres. You will also note that version 17, which was in beta testing when this image was taken, is also available for use. When version 17 is officially released, the "Current" link will point to that version instead, and in this case, version 12 will be removed as it will go end-of-life. Another point is that there are translations of the

manual to various languages linked in the right sidebar, helping make this documentation more accessible.

However, the one key to using the PostgreSQL documentation successfully is to keep an eye on the version. While most search engines will automatically point you to the "Current" documents, this isn't 100% the case and sometimes you will be pointed to an older version of the docs. Similarly, frequently when people link to the documentation from their own blogs and websites, you'll find that they've linked to older versions of the documentation. Figure 17-2 shows the current version of the Indexes page (located here: https://tinyurl.com/3eztm6kk):

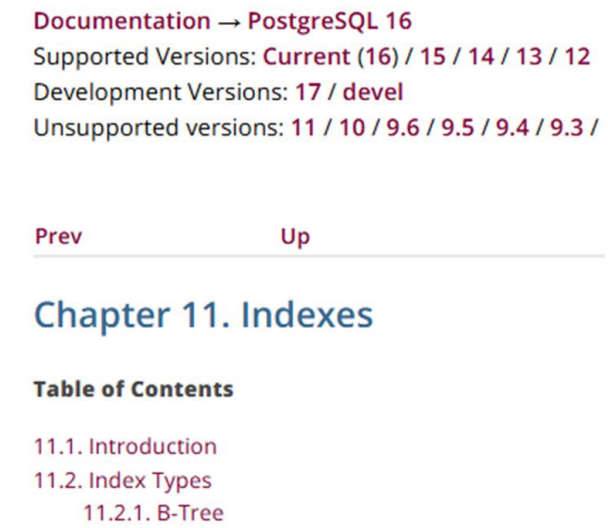

Figure 17-2. *The "Current" documentation on indexes*

While Figure 17-3 shows version 13 of the same landing page:

Figure 17-3. *Version 13 of the documentation on indexes*

While these look almost the same, they're not. The key is to note which version of the documentation you're on. If you look at the top of both Figures 16-2 and 16-3, you'll note that it says what version of the documentation you're currently looking at:

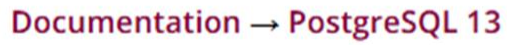

Figure 17-4. *The label that tells you the version of the documentation*

You can also see this if you look at the two URLs:

```
https://www.postgresql.org/docs/13/indexes.html
https://www.postgresql.org/docs/current/indexes.html
```

Notice that version 13 is in the URL, as is "current".

The reason this is important is because you need to ensure that you're looking at the appropriate documentation. Something that works in the latest version of PostgreSQL may not work in version 13. Conversely, things that worked in 13 might not work in the latest version. If you're seeking to learn, you need to ensure that the version you're reading up on applies to the version you're using. Otherwise, you'll have a sometimes-painful process.

While the documentation is extremely good, it's not always laid out in a way that is conducive to how people learn. There also aren't pathways of learning through the documentation. This is why, despite how powerful the documentation is overall, there is a need for other resources to help you in your learning.

Books on PostgreSQL

There are many books on PostgreSQL, including the one you're reading. While no doubt all of them are worthy of your attention, we're going to call out just a few that might stand out from the crowd.

NOTE: None of the links provided here are affiliate links.

The Art of PostgreSQL

The Art of PostgreSQL (https://tinyurl.com/t5veh94e), by Dimitri Fontaine, is one of the most recommended books on PostgreSQL. The book itself focuses primarily on building applications with PostgreSQL , so it's clearly a must read for developers. However, well written SQL is just as important for working in analysis, report writing, or when supporting PostgreSQL, making this an informative read for just about anyone working within PostgreSQL.

Dimitri Fontaine has extensive experience working in development in general and specifically open-source development. He is a major contributor to PostgreSQL itself, so comes from a place of real knowledge and understanding.

Database Administration: The Complete Guide to DBA Practices and Procedures

Database Administration: The complete Guide to DBA Practices and Procedures (https://tinyurl.com/a8h3vxuu) was written by Craig S. Mullins. This isn't a book on PostgreSQL. Instead, it's a book on the job of the Database Administrator. However, since

so many people are doing DBA work, probably without worrying about the title, or have stumbled into the job as "accidental" DBAs (you'll also hear "incidental" or "reluctant" DBAs), knowledge of what's necessary to get the job done makes a big difference in learning PostgreSQL.

Craig Mullins has over 40 years of experience within IT. This is the second edition of the book and is much more modern than the first. Mr. Mullins also wrote several other books on databases and development.

PostgreSQL Query Optimization: The Ultimate Guide to Building Efficient Queries

PostgreSQL Query Optimization: The Ultimate Guide to Building Efficient Queries (https://tinyurl.com/3t952jet), by Henrietta Dombrovskay, Boris Novikov and Anna Bailliekova. The second edition of the book improves on the original, expanding the information available and updating it based on new functionality. You'll get most of what you need to help improve the performance of your PostgreSQL instances.

Henrietta Dombrovskaya has over 40 year's experience within IT and fourteen working with PostgreSQL in particular. Boris Novikov is a database expert and university professor teaching at both undergraduate and post-graduate levels. Anna Bailliekova is a data engineer with more than twelve years of experience in IT.

PostgreSQL Events

Nothing quite beats in-person learning. From the ability to ask specific questions, to the hallway conversations between and around sessions, live, in-person learning has many advantages. There are a number of in-person events and resources that you should know about so that you can take advantage of them as a part of your continued learning within PostgreSQL.

User Groups and Meetups

One of the best in-person resources will always be local user groups, meetups and the like. Not only can you learn from the other attendees, but these are also your peers going through many of the same tribulations you are. You can network with them to make contacts that may help you in your career in other ways as well. Local groups act as a support network and a learning resource.

One of the best places to go for local PostgreSQL resources is right back to postgrsql.org. There is a listing of local user groups under the community page (https://tinyurl.com/3xcaee9n). These are broken down by country and city, so you can track down the one that's right for where you live, or even where you're visiting.

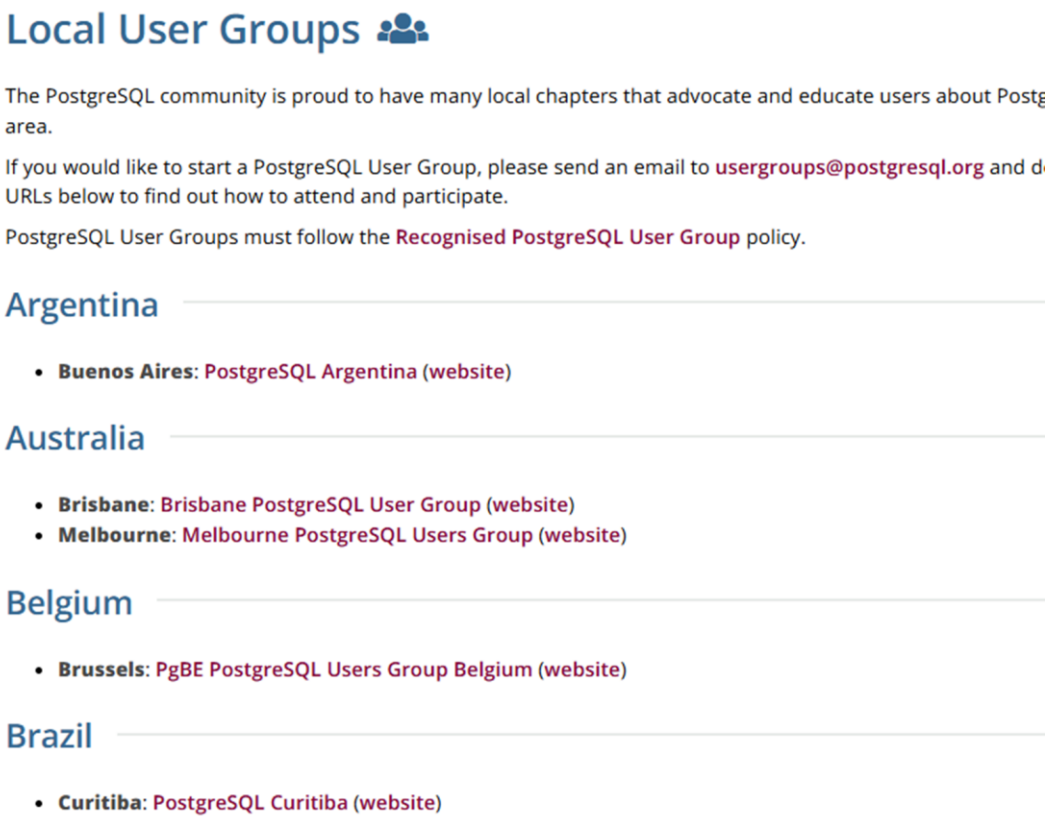

Figure 17-5. *Local User Groups on the PostgreSQL.org page*

Another source for local resources to help you learn PostgreSQL is to look to MeetUp (https://www.meetup.com/). There may be some overlap between what's listed in MeetUp for your area and the Local User Groups listed above. However, a lot of different groups may be teaching PostgreSQL locally. Just as an example, a group in Washington DC called "Data Wranglers DC" meets on various database topics, including PostgreSQL. Another example is the "Berlin Accounting Tech" Meetup that will be hosting an event called "Building an Accounting Ledger with Postgres and Python." In addition to simply searching through MeetUp, there is a designated PostgreSQL page (https://www.meetup.com/pro/postgresql/).

Local Events

A local event is organized much less frequently than a user group. These may be annual or semi-annual and are meant to be a little more regional in scope. There are many different ones organized along different lines. Here are three examples:

PGDay

The PGDay events are organized under the umbrella of the PostgreSQL Community Conference rules (https://tinyurl.com/bfxkmw9n). There are a number of them in different locations around the world. These are typically a one-day event, focused exclusively on PostgreSQL, and often run by one of the User Groups listed on the communities page of PostgreSQL.org.

While these are local and regional events, they draw speakers from around the world. You'll be seeing the kinds of presentations that you'll see at the larger, international events. Upcoming PGDay events will be listed on the Events page of the Community section of PostgreSQL.org (https://tinyurl.com/mtsp2bfx).

DataSaturdays

DataSaturdays are a more generic type of regional event. Instead of a focus exclusively on PostgreSQL, these may include all sorts of other data platforms and other data-focused content. However, they regularly have PostgreSQL content, and some of it from the same speakers that you would see at much larger, international events. These events are

organized locally and run under the rules outlined at the DataSaturdays website (https://datasaturdays.com/). Also, they're not just on Saturday so pay attention to the details and date of the event when you are planning to attend.

SQLSaturday

SQLSaturday started life as a Microsoft Data Platform only event. However, over the years, these one-day, local events, held all over the world, and just like DataSaturdays, held on every day of the week, have grown and expanded to the point where, the SQL in the name does include PostgreSQL sessions. While you won't see as much PostgreSQL content at these events, it is there and growing. SQLSaturday events are organized and run locally under the rules outlined at the SQLSaturday website (https://sqlsaturday.com/).

International Events

Generally, the larger events are almost always multi-day, paid affairs. They are often organized by larger groups of people that are probably not even local to where the event is taking place. These are likely to include one or more days of full-day content, often referred to as pre-cons, meaning pre-conference events. They will then have one or more days with multiple tracks, hosting a very large variety of content. Frequently this is where you'll see, and meet in person, some of the more well-known speakers. These events often draw attendees and speakers from around the world.

PGConf.EU

From a PostgreSQL perspective, PGConf.EU is the conference to attend. PGConf is organized and run by the governing body of PostgreSQL. It is considered not merely a premier event in general, but the showcase for PostgreSQL.

The .EU events are held in various locations around Europe. The 2023 event was in Malta. The 2024 event was in Athens, Greece. Upcoming events will be held in other locations.

In terms of in-person events, this is, without a doubt, the single best resource for PostgreSQL users (we'll talk more about developer resources at the Hackers Conference

below). That's best in terms of the number and quality of sessions, as well as the people presenting them. That's not to say others are not good, especially when you consider they draw from a lot of the same speakers. However, PGConf.EU is very well established within the PostgreSQL community and draws from a larger speaker pool.

PGConf.NYC

Run by the US PostgreSQL organization, PGConf.NYC is another event that draws on both an international attendance and an international pool of speakers. While not as large as PGConf.EU, it's still one of the larger PostgreSQL specific in-person events you could attend.

PGConf.Dev

The PGConf.Dev, formerly PGCon, conference is specifically focused on those individuals who are contributing to PostgreSQL development. As such, it doesn't offer as many general knowledge sessions as many of the other conferences on this list. Instead, it is focused on how to contribute to the PostgreSQL project itself. If you are interested in moving into this aspect of PostgreSQL, or want to get deeper into the PostgreSQL internals, this is probably the single most important conference. The event used to be held exclusively in Ottawa, Canada, but moved to Vancouver in 2024, and will be held in Montreal in 2025.

Postgres Conference

Postgres Conference is a privately run non-profit event focused on training and networking around Postgres technology. The event itself has sessions lead by acknowledged industry experts. It very much focuses on collaboration between the industry and the end-users. In addition to the in-person events, they regularly host online webinars as well.

PASS Data Community Summit

Only within the last two years has the PASS Data Community Summit embraced PostgreSQL as a fundamental aspect of the event. However, this large, in-person event,

usually held in Seattle, WA, is quickly becoming an established destination for PostgreSQL learning. There have been dedicated tracks and pre-cons available on PostgreSQL, with sessions by many of the same people you would see at PGConf.EU or Postgres Conference.

Developer Conferences

Everything mentioned so far has either been conferences dedicated to PostgreSQL, or conferences dedicated to databases and data management. However, there is a lot of PostgreSQL learning available at a whole host of large scale, international events such as:

- FOSDEM
- THAT Conference
- ScalePyCon

Online Resources

To say that there are a lot of resources online for PostgreSQL is probably the understatement of the century, and the century is young. The trick then becomes finding those resources that are generally respected and are worth your time pursuing.

Aggregations

One of the best places to go to get access to online PostgreSQL material has to be Planet Postgres, which is a part of the PostgreSQL.org (https://planet.postgresql.org/). It mostly consists of blogging, but you'll also see webinars and podcasts go by as well. It's voluntary whether people publish through Planet Postgres, but it's still a great resource.

Another place to go for a broad set of contributions is the PostgreSQL Slack (https://postgresteam.slack.com). There are several channels, but of special interest to someone just starting is the Beginners channel. You will find several different channels focused on various kinds of knowledge. There's also a channel on Events if you're looking for something along those lines.

Simple-Talk is an online magazine run by Redgate Software. It pays authors and goes through a rather rigorous editing process, making it a lot more than a blog. You can find a lot of content focused on PostgreSQL there (https://tinyurl.com/bdu22u3j).

Cooper Press publishes a PostgreSQL email newsletter weekly. This acts as an aggregator of blogs and other online content. You can sign up at their web site (https://tinyurl.com/59dv39zw).

Another place that can be a handy place to get a lot of information quickly is X, formerly Twitter. Specifically, going to the #pghelp hash tag, you'll see excellent questions and answers from a very active community. You can either ask questions and learn that way, or, read through people's questions and answers to learn.

One more source for questions and answers, as well as a broad source of information, would be Stack Exchange, but more specifically, the DBA part of Stack Exchange (https://dba.stackexchange.com/). Here again, you can ask questions, answer questions, or simply search to read about topics that you're attempting to learn.

Podcasts

If you enjoy learning by listening, then podcasts are the way to go.

One of the best podcasts out there is Talking Postgres with Claire Giordano(https://talkingpostgres.com/). She and her guests talk about just about everything there is to know about PostgreSQL. She has a very extensive backlog of episodes, so there's a lot to learn.

Another good podcast to learn PostgreSQL from is postgres.fm by Nikolay Samokhvalov and Michael Christofides (https://postgres.fm/). They cover a very wide set of topics, but they all revolve around PostgreSQL. Postgres.fm also has a very extensive backlog of episodes providing more learning.

One more podcast is Scaling PostgreSQL, hosted by Creson Jamison (https://www.scalingpostgres.com/). This weekly show is a little more focused in its topic, as clearly stated in the name. If you're looking to make your PostgreSQL bigger and faster, then this is your podcast.

Blogs

Trying to name individual blogs that are useful will be a difficult undertaking. Instead, we'll focus on a few blog aggregators. The first, we've already mentioned, is Planet Postgres. That would be the main place to start.

PGSQL Phriday is a monthly, themed, blog aggregator run by Ryan Booz but hosted by different people every month (https://www.pgsqlphriday.com/). A person is assigned a given month, and they get to think up a topic, question, or challenge, that a whole host of bloggers then go and create content for. It's been running for 17 months now, so there's a lot of content to pick from.

A slightly different blog aggregator is the PostgreSQL Person of the Week, published by Andreas Scherbaum (https://postgresql.life/). It's not exactly an aggregation of blog posts. Instead, it's a collection of interviews with various people who are involved with the PostgreSQL community. Many of them, however, do maintain blogs, so it's a way to get to know someone to decide if their blog is worth your time.

Webinars

While there are lots of YouTube channels publishing PostgreSQL content, here we'll focus just a few of the regular webinars associated with PostgreSQL, treating video similar to how we treated blogs.

It's a little unfair to call it just an online seminar, because Posette (formerly CitusCon) is much more of an event (https://tinyurl.com/2z2x9nd2). It covers a very wide range of PostgreSQL content. The most recent one hosted 42 sessions by experienced speakers. You can also watch content from earlier events, representing a large collection of content related to PostgreSQL.

Mentioned earlier, Postgres Conference runs a webinar series annually. The link is the same as the one published earlier. The content comes from a variety of people and covers many topics.

Ryan Booz has been running a regular online seminar series entitled PostgreSQL 101 (https://tinyurl.com/y5bd7ttn) hosted by Redgate Software. Clearly, if you're just getting started with PostgreSQL, this series is targeted at you.

One last online resource is the "Postgres for All" meetup, which functions just like any of the in-person MeetUp events, but instead is online. They provide an interactive platform for those learning and working in PostgreSQL (https://tinyurl.com/4b3xybyd).

Conclusion

The goal of this book is to give you the information you need to get started with PostgreSQL. However, we know that this topic, PostgreSQL, includes a lot of information. We're assuming you're going to need more than just our introduction. That's why we included so many other places to learn in this chapter. Whether you want to just go to your local meetup, travel to a large international conference, or learn online at home, we've covered a lot of that material.

Hopefully, you'll find this book helpful. Best of luck on your journey to learn PostgreSQL.

Continue Your PostgreSQL Journey with Redgate

As part of our commitment to supporting the Postgres community, we're excited to offer two new tools from Redgate that can help you go even further:

- **Redgate pgNow** – Stay on top of your Postgres estate with real-time insights and monitoring with this free tool.
- **Redgate pgCompare** – Effortlessly compare and synchronize Postgres schemas with precision and ease.

Simply scan the QR codes below to explore these tools and see how they can support your next steps with Postgres.

Redgate pgNow

Redgate pgCompare

Index

D

E

F

P

Q

R

S

T

U

V

W

X

Y